

Building Systems for Scalable and Responsible Data Management

The Data-Intelligence Loop

Zhengyi Yang

University of Sydney

zhengyi.yang@sydney.edu.au

Office: Room 350, Building J12

Motivation

My research asks how to build data systems that scale to complex workloads, support reliable AI workflows over data, and inform responsible real-world decision-making.

Complex and dynamic data

Tables, graphs, text, documents, multimedia data, logs, streams, and evolving workloads.

AI data workloads

Retrieval, indexing, evidence selection, reasoning, verification, serving, and cost control.

Real-world applications

ESG, finance, workforce analytics, fairness, climate, health, and scientific discovery.

Database systems provide the foundation for managing, processing, and reasoning over data across these motivations.

Three-Theme Research Agenda

STAGE



Data Processing

Process heterogeneous and dynamic data efficiently.



Data Understanding

Extract structure, semantics, context, and evidence.



Data Intelligence

Support responsible decisions in real domains.

RESEARCH FOCUS

THEME 1

Scalable Data Systems

Graphs, tables, streams, query processing, GPU, out-of-core

THEME 2

Responsible Data Intelligence

GraphRAG, LLM reasoning, AI/ML for DB, tabular learning, data-centric AI


THEME 3

Data for Real-World Applications

ESG, finance, built environment, fairness, climate, health

The agenda links scalable systems, responsible intelligence, and real-world applications through three themes.

The Feedback Loop



BOTTOM-UP

Capabilities accumulate

Scalable systems make complex data usable; data understanding builds retrieval, reasoning, and learning on top.

TOP-DOWN

Requirements return

Real-world domains send requirements back to AI pipelines and system design: provenance, robustness, fairness, validation, and deployability.

Processing enables understanding; understanding enables intelligence; real-world applications feed requirements back into data systems.

PART I

Scalable Data Systems

Data processing is about making complex and dynamic data computationally manageable at scale.

Graph Analytics and Processing

Graphs, hypergraphs, temporal graphs, matching, centrality, decomposition.

Database Systems & Query Processing

Query processing, joins, JIT/runtime, OLTP kernels, storage, transactions.

Parallel and Hardware-Aware Processing

GPU, FPGA, RISC-V, UMA, streaming, parallelism, memory- and I/O-aware execution.

Graph Analytics at Scale

Complex relational structures are no longer only graphs; they include higher-order and temporal interactions that make naive enumeration infeasible.

REPRESENTATIVE TOPICS

- **Higher-order** graph and hypergraph analytics
- **Pattern matching** and substructure search
- **Connectivity** and path analytics
- **Temporal** and evolving graph analytics

Turn structural properties into indexes, pruning rules, and execution strategies that avoid exhaustive search.

SELECTED WORK

COMPLETED

Nucleus decomposition

SIGMOD 2026

Diversified graph pattern matching

VLDB 2026

Hypergraph pattern matching

ICDE 2026

ICDE 2023

Hypergraph decomposition

SIGMOD 2025

VLDBJ 2025

CIKM 2024

Constrained path analytics

ICDE 2025

Clique cover and dense subgraph

WWW 2025

Temporal graph ranking and analytics

WWW 2024

WWWJ 2024

UNDER SUBMISSION

Higher-order graph analytics

SPRINGER MONOGRAPH, IN PRESS

Nucleus and nuclear decomposition

SIGMOD 2027

VLDB 2027

Hypergraph core decomposition and maintenance

ICDE 2027

VLDBJ

Reachability indexing and querying

EDBT 2026 UNDER REVISION

Database Systems & Query Processing


Database performance depends not only on algorithms but also on how query logic interacts with compilation, runtime execution, storage, and deployment constraints.

REPRESENTATIVE TOPICS

- **Transactional** engine design
- **Query processing**, joins, and runtime optimization
- **JIT / MLIR** compilation for database execution
- **Storage**, concurrency, and recovery

Systems optimization should make database engines faster while keeping them maintainable and deployable.

EXECUTION PIPELINE



SELECTED WORK

PhoebeDB: High-Performance Disk-Based RDBMS Kernel for OLTP

EDBT 2025 INDUSTRY INDUSTRY COLLABORATION TARGET: INDUSTRY-READY

- **Industry collaboration** toward an **enterprise-oriented**, industry-ready RDBMS kernel.
- **PostgreSQL-compatible OLTP** on **commodity hardware**, reducing migration and deployment risk.
- **Full-stack kernel design**: temperature-aware storage, coroutine runtime, smart scheduling, transaction management, WAL, and snapshots.
- Reports **13.7M tpmC** and **30M tpm** on **TPC-C** using one machine, with **27x** improvement over PostgreSQL.

● COMPLETED

Batch-dynamic high-dim kNN join

ADMA 2024 BEST STUDENT PAPER

● UNDER SUBMISSION

Fully dynamic high-dim kNN joins

VLDBJ

Parallel and Hardware-Aware Processing

At billion scale, performance depends on coordinating distributed execution, parallelism, memory hierarchy, storage, and accelerators.

REPRESENTATIVE TOPICS

- **Distributed** and out-of-core processing
- **Parallel** runtime and adaptive scheduling
- **GPU-/FPGA-accelerated** indexing and analytics
- **Emerging hardware** and scalable execution

Scale by co-designing algorithms, runtimes, and hardware-aware execution.

EXECUTION DIMENSIONS

Distributed execution

Memory / Out-of-core

Parallel runtimes

GPU / FPGA / RISC-V

SELECTED WORK

COMPLETED

Out-of-core graph processing system

SIGMOD 2026

GPU-accelerated vector index construction

SIGMOD 2026

GPU-accelerated range top-k query

HPDC 2026

Database benchmarking on RISC-V

VLDB ADMS 2025

UNDER SUBMISSION

FPGA-accelerated graph random walks

SIGMOD 2027

SIMD set-intersection benchmarking

SIGMOD 2027

Out-of-core vector search and indexing

EUROSYS 2027

Unified-memory DataFrame processing

VLDB 2026 DEMO

PART II

Responsible Data Intelligence

Data understanding is about turning processed data into structure, semantics, context, and reliable evidence for AI systems.

Data Systems for Reliable LLMs

Retrieval, evidence selection, reasoning, agent memory, and context management.

AI/ML for Data Systems

Cost estimation, tuning, temporal analytics, and costly system decisions.

Data-Centric AI

Tables, labels, condensation, reasoning, quality, and data-efficient learning.

Data Systems for Reliable LLMs


Reliable LLM applications over data need systems for retrieval, evidence organization, verification, reasoning orchestration, and serving, not just fluent generation.

REPRESENTATIVE TOPICS

- **Data-to-structure** construction
- **Evidence retrieval** and selection
- **Reasoning** and verification
- **Serving, memory**, and context management

Reliable LLMs create data-management workloads: indexing, planning, provenance, cost control, and evidence quality.

RETRIEVAL-TO-REASONING PIPELINE



SELECTED WORK

COMPLETED

Adaptive GraphRAG retrieval

KDEXLLM 2026 BEST PAPER

Graph exploration for data understanding

SIGMOD DEMO 2025

UNDER SUBMISSION

Graph CoT reasoning and serving

TOIS

Cost-efficient KG evidence construction

TKDE

Semantic-unit retrieval for HyperRAG

EMNLP 2026

Long-term memory for LLM agents

EMNLP 2026

AI/ML for Databases and Data Systems

This line of work learns to approximate, tune, and adapt costly decisions in data systems while respecting query, graph, and temporal semantics.

REPRESENTATIVE TOPICS

- **Learning** for query optimization
- **Graph Neural Network** (GNN)
- **Representation learning** over structured data
- **Deployment-aware** learned components

AI/ML is useful when learning acts as a system component with clear semantics, error behavior, and operational cost.

Query Optimization

Cost estimation and optimization signals for expensive database decisions.

Graph Analytics

Temporal centrality and graph signals where exact computation is costly.

SELECTED WORK

● COMPLETED

Learning-based temporal centrality

WWW 2026

WWW 2024

Adaptive parallelism tuning for stream processing

ICDE 2025

Heterogeneous graph learning

CIKM 2025

● UNDER SUBMISSION

Learned query cost estimation

ICDE 2026

Event-centric HIN learning

ICDM 2026

Learning-based attributed community detection

TKDD

Data-Centric AI

Data-centric AI studies how tables, labels, schemas, and data quality shape model learning, reasoning, efficiency, and reliability.

REPRESENTATIVE TOPICS

- **Structured data** representation and reasoning
- **Data condensation** and efficiency
- **Label quality** and aggregation
- **Data quality** for AI workflows

Data-centric AI brings classic data quality questions back into the core of model performance and reasoning reliability.

DATA PROBLEMS

Tables

Representation, schema, values, and reasoning.

Labels

Crowdsourcing, quality, and aggregation.

Efficiency

Condensation and data-efficient learning.

SELECTED WORK

● COMPLETED

Tabular data condensation

ICDE 2026

Structured table reasoning

WWWJ 2025

● UNDER SUBMISSION

LLM-based tabular reasoning

EMNLP 2026

Label aggregation in crowdsourcing

ICDE 2026

PART III

Data for Real-World Applications

Data intelligence is about using processed and interpreted data to support responsible decisions in real-world domains.

ESG and Finance

Disclosure, standards, fraud, interpretable analytics.

Built Environment

Smart construction, careers, leadership, guidance.

Fairness and Diversity

Representation, bias, evaluation, context.

Science and Society

Climate, health, materials, validation.

ESG and Financial Intelligence

ESG (Environmental, Social, and Governance) and finance turn messy documents, reports, and behavioral traces into decisions that require traceability and interpretation.

REPRESENTATIVE TOPICS

- **Document** and report intelligence
- **Knowledge-grounded** evidence management
- **Financial behavior** and risk analytics
- **Interpretable** decision support

The database problem is evidence governance: linking claims, standards, provenance, and models into auditable workflows.

ESG

Standards, reports, provenance, auditability.

Finance

Noisy behavior, robust signals, fraud detection.

SELECTED WORK

COMPLETED

Agent-based ESG report analysis

ICSA 2026

Finance analysis and fraud detection

BIGCOMP 2026

ADMA 2025

UNDER SUBMISSION

ESG data management survey

THE INNOVATION

Standard-aware ESG disclosure analysis

ICDM DEMO 2026

ESG knowledge graph construction

ICKG 2026

AI-Empowered ESG Data Analytics

ACTIVE ARC PROJECT

AI-Empowered ESG Data Analytics in Responsible Investment

The project leverages advanced AI and data management to improve ESG risk assessment, performance forecasting, and report generation.

BACKGROUND

ESG links environmental impact, social responsibility, and governance quality to investment, compliance, and accountability.

SIGNIFICANCE

Reliable data infrastructure makes ESG claims traceable, comparable, and actionable for responsible investment.

SCHEME

Early Career Industry Fellowships

LEAD

Dr Zhengyi Yang

FOR

4605 - Data Management and Data Science

ESG analysis is a data-management problem before it is an AI problem: fragmented standards, disclosures, market signals, enterprise evidence, risk signals, forecasts, and reports must become timely, objective, and auditable decision support.

PROJECT AIMS	Assess	Forecast	Generate
	Improve ESG risk analysis from complex, mixed-source data.	Surface performance signals for responsible investment.	Produce transparent reports backed by auditable evidence.

NEED -> OUTPUT -> DATA -> SYSTEM

National need

Reliable ESG insights support sustainable investment, corporate responsibility, and regulatory compliance.

Research output

More accurate, accessible, timely, and objective ESG analysis for investors and companies.

Data science challenges

Availability, quality, integration, traceability, timeliness, privacy, and cross-standard semantics.

Systems agenda

ESG knowledge layer, claim-level provenance, retrieval/reasoning guardrails, and workflow evaluation.

GOVERNED ESG DECISION LOOP

Report

Transform standards, disclosures, and evidence into structured ESG claims.



Review

Check completeness, consistency, standard alignment, and evidence quality.



Respond

Support risk, compliance, investment, and operational decisions with audit trails.

Built Environment Intelligence

Built-environment and smart-construction applications transform fragmented professional, project, and regulatory traces into interpretable evidence for career, compliance, sustainability, and organizational decisions.

REPRESENTATIVE TOPICS

- **Built-environment** decision support
- **Workforce** and career intelligence
- **Compliance** and sustainability analytics
- **Heterogeneous** domain data integration

Real-world decision support often begins as entity resolution, evidence extraction, and schema design over noisy human-centered data.

DATA TO DECISION FLOW



SELECTED WORK

● COMPLETED

Construction career guidance

BIGCOMP 2026 ISARC 2026

Built-environment GraphRAG systems

BUILDINGS 2026

Workforce leadership and progression analysis

BUILDINGS 2026 CRIOCM 2026 IJCM 2025

● UNDER SUBMISSION

Australian built-environment AI co-pilot

AUTOMATION IN CONSTRUCTION

Compliance and sustainability intelligence

ARC LINKAGE PROJECT

Fairness and Diversity Intelligence

Fairness work requires careful data modeling because representation, missingness, and measurement choices shape downstream conclusions.

REPRESENTATIVE TOPICS

- **Representation** and measurement
- **Fairness** in decision support
- **Diversity** and progression analytics
- **Evaluation** of model behavior

Fairness is not only a model metric; it is also a data-management problem of measurement, provenance, and context.

FAIRNESS QUESTIONS

Representation

Who is visible in the data?

Measurement

What proxies shape conclusions?

Decision support

How are outputs used?

Evaluation

What is fairness beyond accuracy?

SELECTED WORK

COMPLETED

Gender representation and career progression

JCEM 2026

IJCM 2025

IJCM 2025

LLM fairness and bias evaluation

FAIML 2026

AI 2025

Scientific and Societal Intelligence

Scientific and societal applications need AI pipelines that respect domain constraints and produce outputs that can be validated.

REPRESENTATIVE TOPICS

- **Domain-aware** scientific modeling
- **Health** and conversation intelligence
- **Scientific discovery** workflows
- **Validation** and trustworthy outputs

Domain-facing intelligence needs data systems that make heterogeneous evidence scalable, checkable, and reusable.

DOMAIN PILLARS

Climate

Heterogeneous evidence and domain-aligned modeling.

Health

Conversation data, semantics, and responsible interpretation.

Science

Materials discovery and validation-centered AI pipelines.

SELECTED WORK

● UNDER SUBMISSION

Weather and climate modeling

NEURIPS 2026

Health conversation data intelligence

ADMA 2026

AI-assisted materials discovery

MULTIPLE MANUSCRIPTS

What This Agenda Enables

A database systems agenda for AI-era decision-making.

The common thread is to turn complex data into scalable computation, reliable understanding, and responsible decision support in real domains.

SYSTEMS

AI PIPELINES

APPLICATIONS

CAPABILITY

Scalable data substrates

Algorithms, indexes, runtimes, and hardware-aware execution for complex workloads.

RELIABILITY

AI-facing data understanding

Retrieval, reasoning, provenance, and verification over structured and unstructured data.

IMPACT

Domain-grounded intelligence

Real-world settings feed back requirements for fairness, robustness, auditability, and deployment.

The broader contribution is to connect database systems, AI pipelines, and real-world data deployment through one coherent feedback loop.

Research Roadmap

DIRECTION 1

Data Systems for Reliable LLMs

- **RAG-aware** indexing and retrieval planning
- **Evidence-aware** query processing and verification
- **Cost-based** optimization for reasoning pipelines
- **Provenance** and audit trails for LLM outputs

DIRECTION 2

Scalable Processing for Complex Data

- **Distributed** graph, hypergraph, and temporal analytics
- **Hardware-aware** vector and graph indexing
- **Adaptive parallel** runtimes for evolving workloads
- **Emerging architectures** for scalable execution

DIRECTION 3

Responsible Data Intelligence

- **ESG** and financial data intelligence
- **Built environment** and smart construction
- **Fairness**, diversity, and accountable data use
- **Climate, health**, and scientific data intelligence

INDEXING

OPTIMIZATION

RUNTIME SYSTEMS

DISTRIBUTED PROCESSING

DATA QUALITY

PROVENANCE

BENCHMARKS

Collaboration & Discussion

RESEARCH SUMMARY

Scalable Data Systems

Graph and hypergraph analytics, database runtime, streaming, and hardware-aware processing.

Responsible Data Intelligence

GraphRAG, LLM retrieval and reasoning, AI4DB, tabular learning, and data-centric AI.

Data for Real-World Applications

ESG, finance, built environment, fairness, climate, and health.

POSSIBLE COLLABORATIONS

RDBMS and HTAP systems

OLTP/HTAP engines, query execution, isolation and consistency, runtime optimization, and provenance.

Data engineering for AI and science

ML-in-DB, multicore/cloud execution, scientific data management, caching, and benchmarking.

Graph data management at scale

Higher-order and temporal graphs, subgraph matching, dynamic graph processing, and GraphRAG indexing.

Thank you!

zhengyi.yang@sydney.edu.au
Room 350, Building J12